

Database Design

DDL Commands

```
CREATE TABLE User (  
    userID INT PRIMARY KEY,  
    userName VARCHAR(255),  
    email VARCHAR(255),  
    universityID INT,  
    FOREIGN KEY (universityID) REFERENCES University(universityID)  
);  
  
CREATE TABLE Family (  
    familyID INT PRIMARY KEY,  
    leaderID INT,  
    accessType VARCHAR(255),  
    serviceName VARCHAR(255),  
    FOREIGN KEY (leaderID) REFERENCES User(userID),  
    FOREIGN KEY (serviceName) REFERENCES SubscriptionService(serviceName)  
);  
  
CREATE TABLE SubscriptionService (  
    serviceName VARCHAR(255) PRIMARY KEY,  
    price DECIMAL(10,2),  
    maxMembers INT  
);  
  
CREATE TABLE University (  
    universityID INT PRIMARY KEY,  
    universityName VARCHAR(255),  
    city VARCHAR(255)  
);  
  
CREATE TABLE Membership (  
    memberID INT,  
    familyID INT,  
    memberStatus VARCHAR(255),  
    PRIMARY KEY (memberID, familyID),  
    FOREIGN KEY (memberID) REFERENCES User(userID),  
    FOREIGN KEY (familyID) REFERENCES Family(familyID)  
);  
  
CREATE TABLE BankAccount (  
    accountName VARCHAR(255),  
    platform VARCHAR(255),  
    userID INT,  
    PRIMARY KEY (accountName, platform),  
    FOREIGN KEY (userID) REFERENCES User(userID)  
);
```

```
CREATE TABLE Payment (  
    paymentID INT PRIMARY KEY,  
    payerID INT,  
    recipientID INT,  
    amount DECIMAL(10,2),  
    paid BIT,  
    deadline DATE,  
    FOREIGN KEY (payerID) REFERENCES User(userID),  
    FOREIGN KEY (recipientID) REFERENCES User(userID)  
);
```

Advanced Queries

Query 1

Our first query returns the number of pending user invitations per university. We first join the **University** table with the **User** table, then group by the **universityID**. We order the output by descending number of users.

Code

```
select un.universityName, count(*) as numUsers  
from University un  
natural join User us  
group by un.universityID  
order by numUsers desc;
```

Result

Python

```

-----+
| -> Sort: numPending DESC (actual time=1.440..1.440 rows=6 loops=1)
|   -> Table scan on <temporary> (actual time=0.001..0.002 rows=6 loops=1)
|     -> Aggregate using temporary table (actual time=1.424..1.425 rows=6 loops=1)
|       -> Nested loop inner join (cost=325.35 rows=369) (actual time=0.097..1.100 rows=369 loops=1)
|         -> Nested loop inner join (cost=196.20 rows=369) (actual time=0.090..0.799 rows=369 loops=1)
|           -> Index lookup on m using memId (memberStatus='Pending') (cost=67.05 rows=369) (actual time=0.046..0.157 rows=369 loops=1)
|             -> Filter: (us.universityID is not null) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=369)
|               -> Single-row index lookup on us using PRIMARY (userID=m.memberID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=369)
|             -> Single-row index lookup on un using PRIMARY (universityID=us.universityID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=369)
|
|-----+
1 row in set (0.00 sec)

```

We see a speedup in the "Filter" command that addresses `m.memberStatus = "Pending"`, and an overall speedup in the whole query, so we can conclude that this index is a good optimization to add for this specific query.

Index 2

For our second index design, we added an index on `University.universityName` because we use it in the grouping step of our query.

```
alter table University add index uniName(universityName);
```

The output of our analyze query follows:

```

| -> Sort: numPending DESC (actual time=2.886..2.886 rows=6 loops=1)
|   -> Table scan on <temporary> (actual time=0.001..0.002 rows=6 loops=1)
|     -> Aggregate using temporary table (actual time=2.867..2.868 rows=6 loops=1)
|       -> Nested loop inner join (cost=683.50 rows=400) (actual time=0.061..2.541 rows=369 loops=1)
|         -> Nested loop inner join (cost=543.50 rows=400) (actual time=0.055..2.211 rows=369 loops=1)
|           -> Filter: (m.memberStatus = 'Pending') (cost=403.50 rows=400) (actual time=0.047..1.515 rows=369 loops=1)
|             -> Table scan on m (cost=403.50 rows=4000) (actual time=0.041..1.113 rows=4000 loops=1)
|               -> Filter: (us.universityID is not null) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=369)
|                 -> Single-row index lookup on us using PRIMARY (userID=m.memberID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=369)
|               -> Single-row index lookup on un using PRIMARY (universityID=us.universityID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=369)
|
|-----+
1 row in set (0.01 sec)

```

Query 2

Our second query returns the number of accepted users for every university and subscription service. We begin by joining the `User`, `Membership`, `Family`, `SubscriptionService`, and `University` tables. We then filter by accepted members, and group by both `universityID` and `serviceName`. We order the output by the `universityName` in ascending order and aggregated `numUsers` in descending order. Finally, we limit the length of the output to 15 rows for visualization.

Code

```

select un.universityName, ss.serviceName, count(*) as numUsers
from User us
join Membership m
  on m.memberID = us.userID
natural join Family f
natural join SubscriptionService ss
natural join University un
where m.memberStatus = "Accepted"

```

```
group by un.universityID, ss.serviceName
order by un.universityName asc, numUsers desc
limit 15;
```

Result

```
with pool.connect() as con:
    statement = '''
    select un.universityName, ss.serviceName, count(*) as numUsers
    from User us
    join Membership m
        on m.memberID = us.userID
    natural join Family f
    natural join SubscriptionService ss
    natural join University un
    where m.memberStatus = "Accepted"
    group by un.universityID, ss.serviceName
    order by un.universityName asc, numUsers desc
    limit 15;
    '''
    result = con.execute(statement).fetchall()
    for row in result:
        print(row)
```

[118] ✓ 0.1s

Python

```
... ('Harper College', 'Netflix', 135)
    ('Harper College', 'Spotify', 131)
    ('Harper College', 'Amazon Prime Video', 126)
    ('Harper College', 'Hulu', 99)
    ('Harper College', 'Disney+', 90)
    ('Harvard University', 'Netflix', 107)
    ('Harvard University', 'Disney+', 106)
    ('Harvard University', 'Amazon Prime Video', 100)
    ('Harvard University', 'Spotify', 100)
    ('Harvard University', 'Hulu', 93)
    ('Massachusetts Institute of Technology', 'Spotify', 112)
    ('Massachusetts Institute of Technology', 'Amazon Prime Video', 104)
    ('Massachusetts Institute of Technology', 'Netflix', 103)
    ('Massachusetts Institute of Technology', 'Disney+', 93)
    ('Massachusetts Institute of Technology', 'Hulu', 93)
```

Indexing

For the second query, we tried another three index designs.